



Computer Networks

Wenzhong Li
Nanjing University



Chapter 5. End-to-End Protocols

- Transport Services and Mechanisms
- User Datagram Protocol (UDP)
- Transmission Control Protocol (TCP)
- TCP Congestion Control
- Real-time Transport Protocol (RTP)
- Session Initiation Protocol (SIP)
- Real Time Streaming Protocol (RTSP)



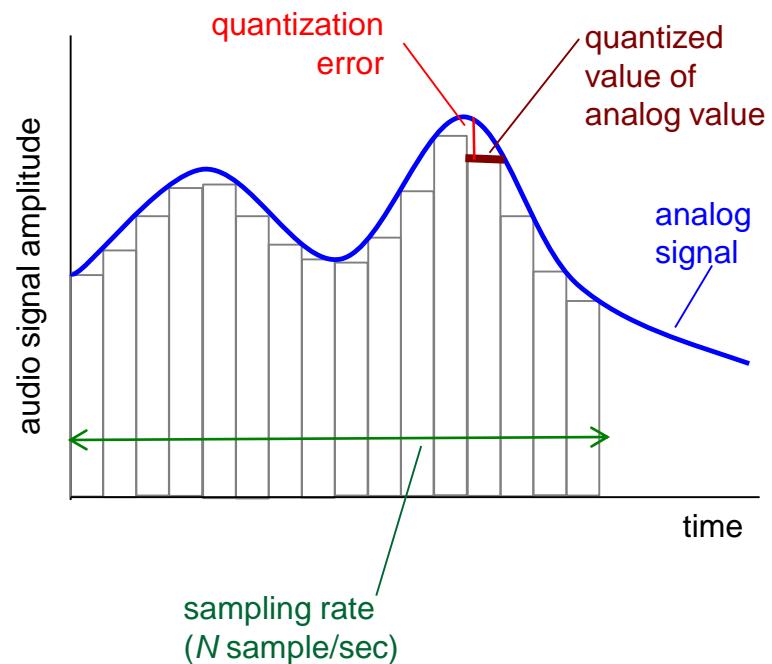
Multimedia Networking





Multimedia: audio

- ❖ Analog audio signal sampled at constant rate
 - telephone: 8,000 samples/sec
 - CD music: 44,100 samples/sec
- ❖ Each sample quantized, i.e., rounded
 - e.g., $2^8=256$ possible quantized values
- ❖ Example rates
 - CD: 1.411 Mbps
 - MP3: 96, 128, 160 kbps
 - Internet telephony: 5.3 kbps and up



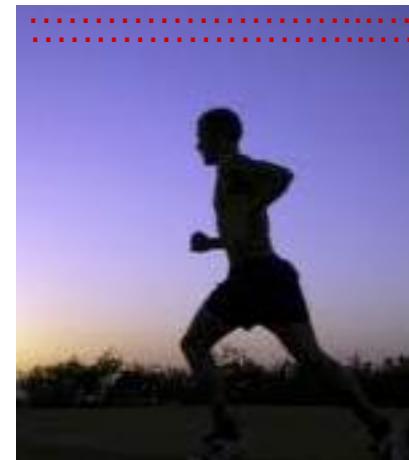


Multimedia: video



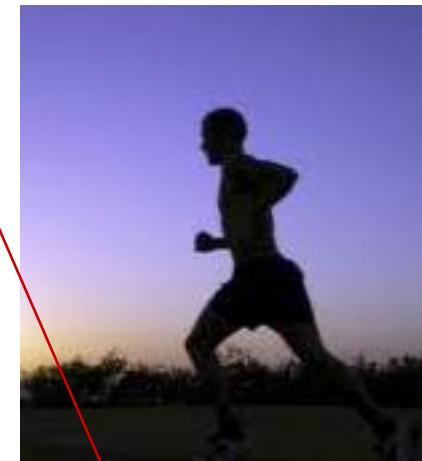
- ❖ Video: sequence of images displayed at constant rate
 - e.g. 24 images/sec
- ❖ Digital image: array of pixels
 - each pixel represented by bits
- ❖ Coding: use redundancy *within* and *between* images to decrease # bits used to encode image
 - spatial (within image)
 - temporal (from one image to next)
- Examples:
 - MPEG 1 (CD-ROM) 1.5 Mbps
 - MPEG2 (DVD) 3-6 Mbps
 - MPEG4 (often used in Internet, < 1 Mbps)

spatial coding example: instead of sending N values of same color (all purple), send only two values: color value (*purple*) and number of repeated values (N)



frame *i*

temporal coding example: instead of sending complete frame at $i+1$, send only differences from frame *i*



frame *i+1*

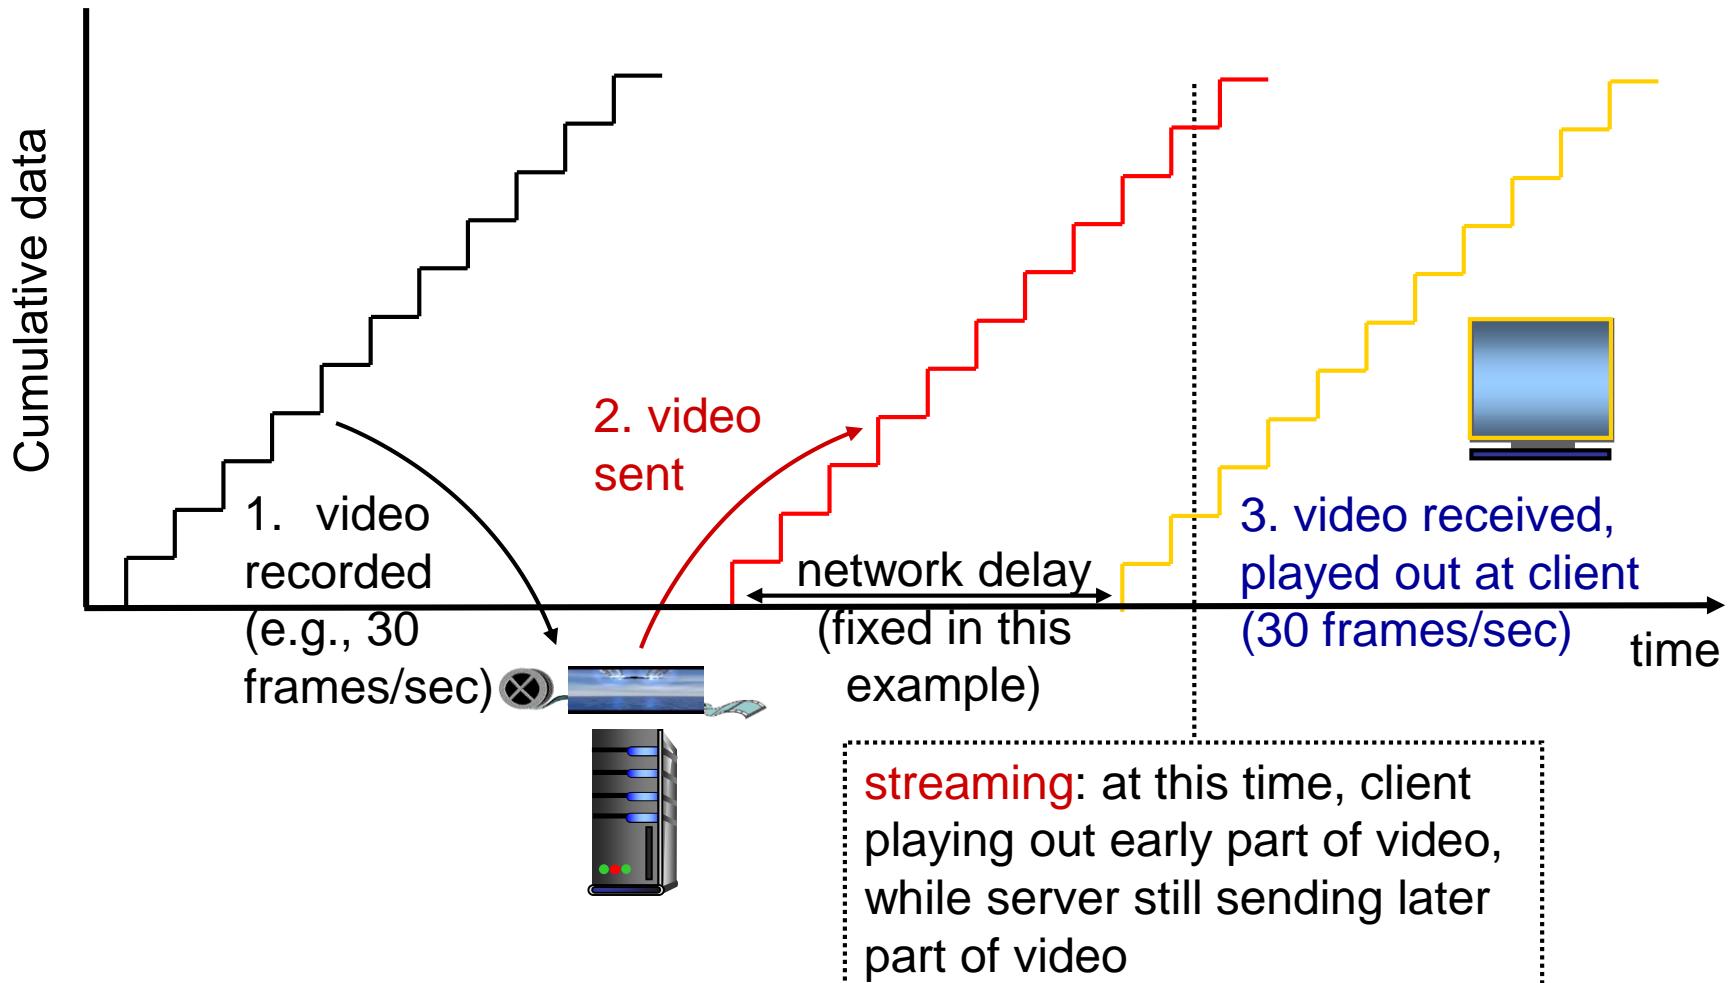


Three application types

- ❖ *Streaming, stored* audio, video
 - *streaming*: can begin playout before downloading entire file
 - *stored (at server)*: can transmit faster than audio/video will be rendered (implies storing/buffering at client)
 - e.g., YouTube, Netflix, Hulu
- ❖ *Conversational* voice/video over IP
 - interactive nature of human-to-human conversation limits delay tolerance
 - e.g., Skype
- ❖ *Streaming live* audio, video
 - e.g., live sporting event (futbol)



Streaming stored video:



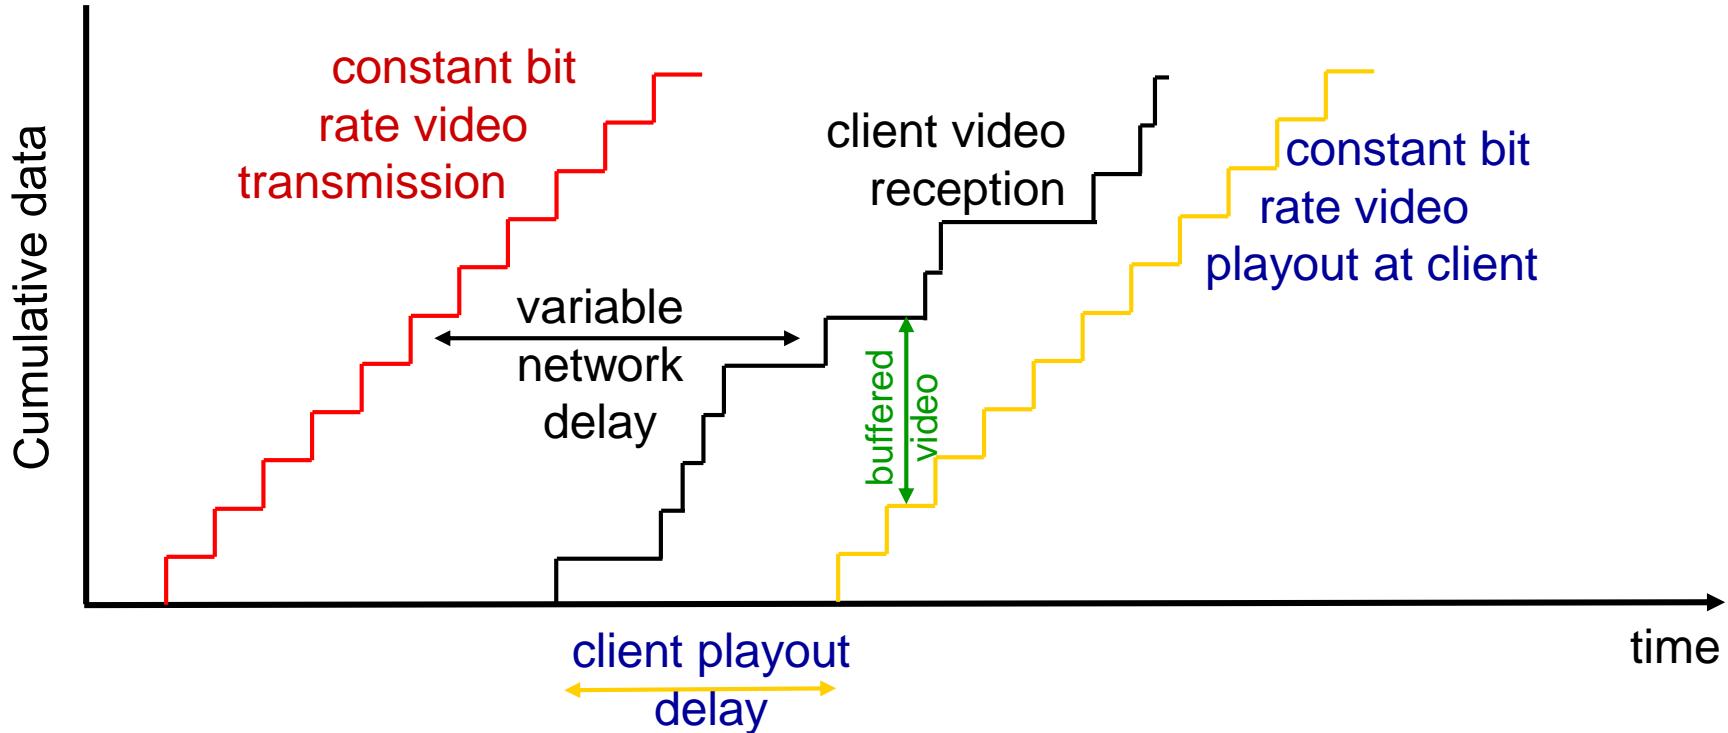


Streaming stored video: challenges

- ❖ *Continuous playout constraint*: once client playout begins, playback must match original timing
 - ... but *network delays are variable* (jitter), so will need *client-side buffer* to match playout requirements
- ❖ Other challenges:
 - client interactivity: pause, fast-forward, rewind, jump through video
 - video packets may be lost, retransmitted



Streaming stored video: revisited



- ❖ *client-side buffering and playout delay:* compensate for network-added delay, delay jitter



Streaming multimedia: UDP

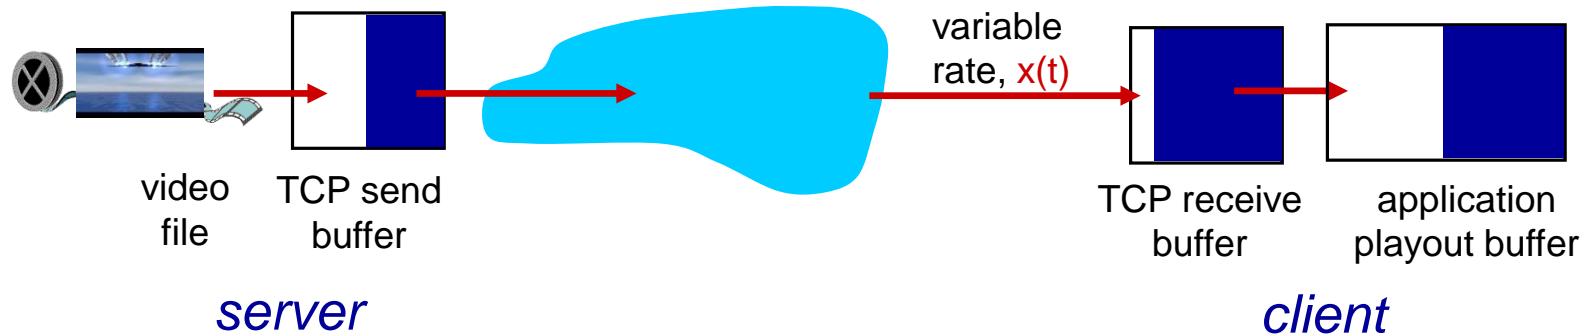


- ❖ Server sends at rate appropriate for client
 - often: send rate = encoding rate = constant rate
 - transmission rate can be oblivious to congestion levels
- ❖ short playout delay (2-5 seconds) to remove network jitter
- ❖ error recovery: application-level, time permitting
- ❖ RTP [RFC 2326]: multimedia payload types
- ❖ UDP may *not* go through firewalls



Streaming multimedia: HTTP

- multimedia file retrieved via HTTP GET
- send at maximum possible rate under TCP



- fill rate fluctuates due to TCP congestion control, retransmissions (in-order delivery)
- larger playout delay: smooth TCP delivery rate
- HTTP/TCP passes more easily through firewalls



Streaming multimedia: DASH

- ❖ *DASH: Dynamic, Adaptive Streaming over HTTP*
- ❖ *server:*
 - divides video file into multiple chunks
 - each chunk stored, encoded at different rates
 - *manifest file:* provides URLs for different chunks
- ❖ *client:*
 - periodically measures server-to-client bandwidth
 - consulting manifest, requests one chunk at a time
 - chooses maximum coding rate sustainable given current bandwidth
 - can choose different coding rates at different points in time (depending on available bandwidth at time)



Streaming multimedia: DASH



- *DASH: Dynamic, Adaptive Streaming over HTTP*
- "*intelligence*" at client: client determines
 - *when* to request chunk (so that buffer starvation, or overflow does not occur)
 - *what encoding rate* to request (higher quality when more bandwidth available)
 - *where* to request chunk (can request from URL server that is "close" to client or has high available bandwidth)



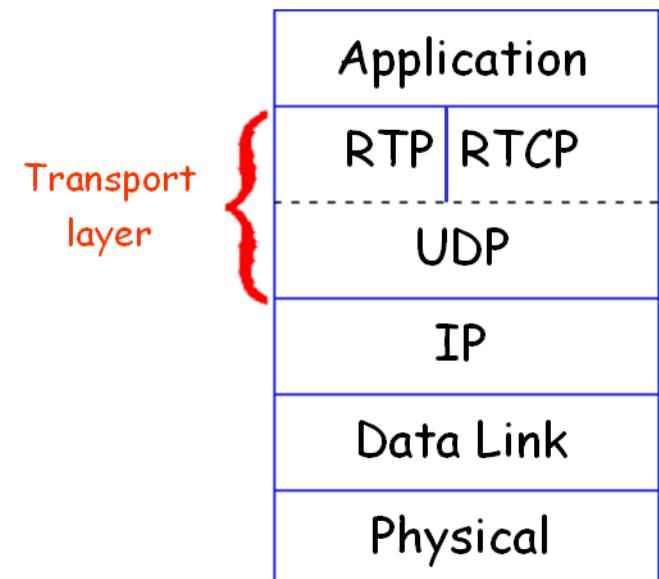
Real-time Transport Protocol (RTP)

■ RFC 3550

- Built upon UDP, i.e. RTP packets encapsulated in UDP segments
- Specifies packet structure for packets carrying audio, video data

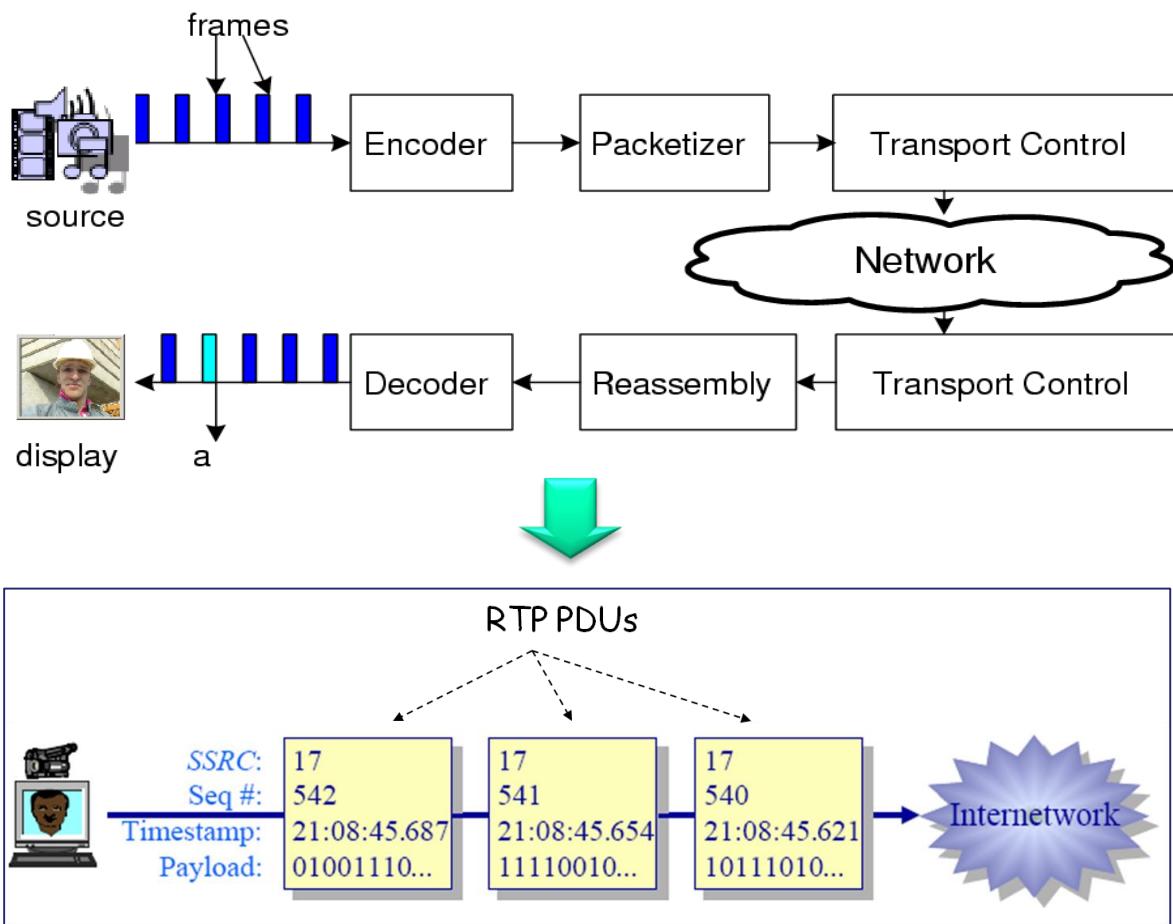
■ RTP packet provides

- Payload type identification
- Packet sequence numbering
- Time stamping



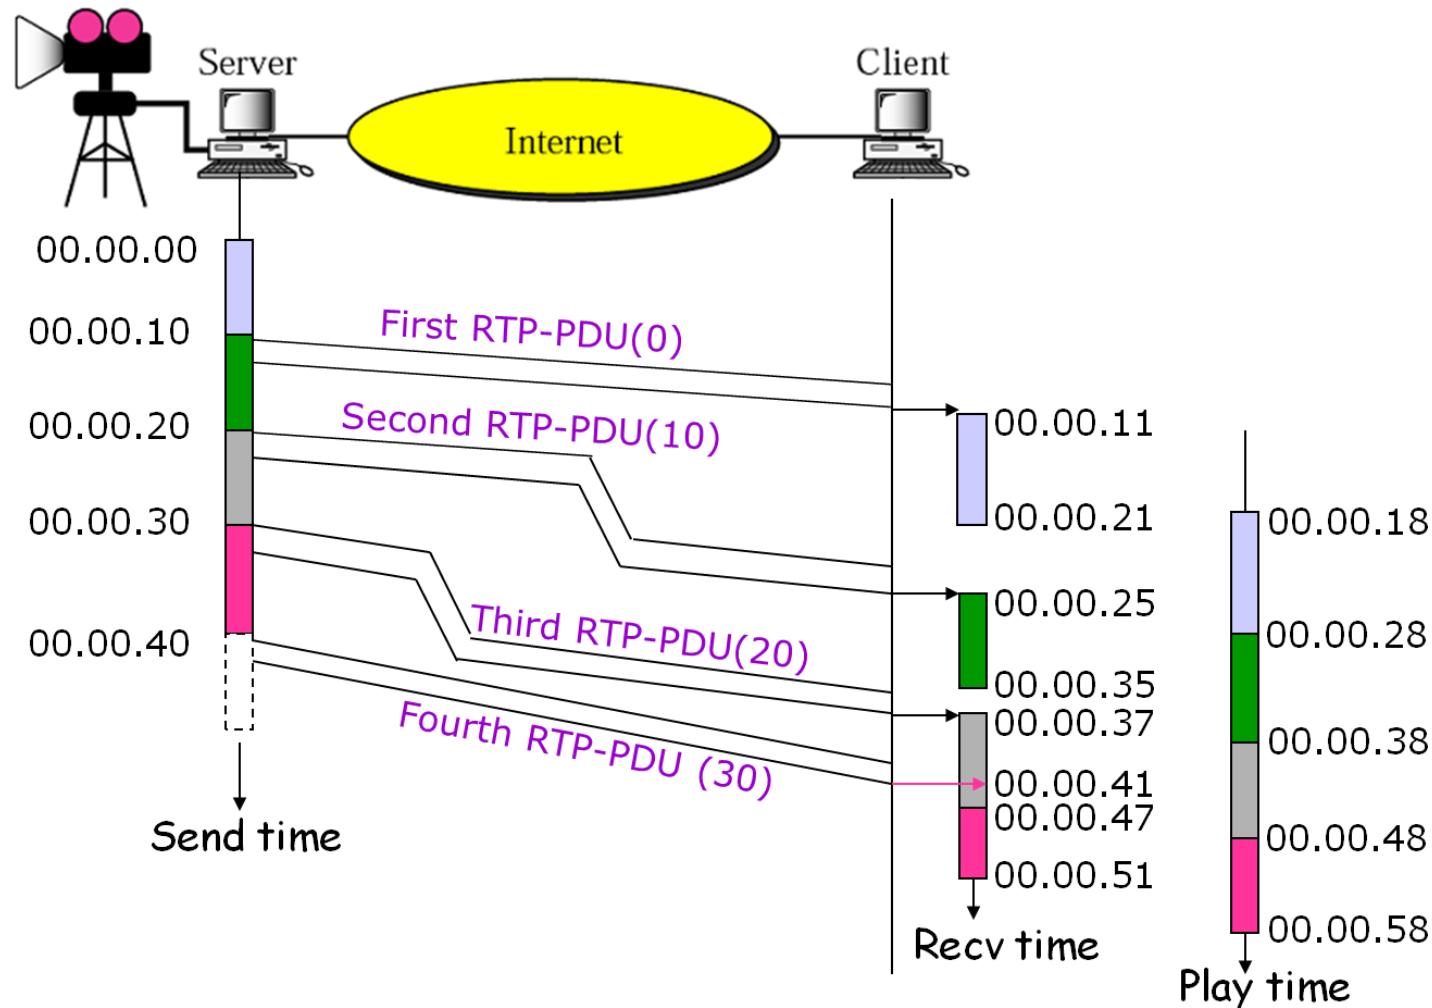


Real-time Streaming





Real-time Streaming



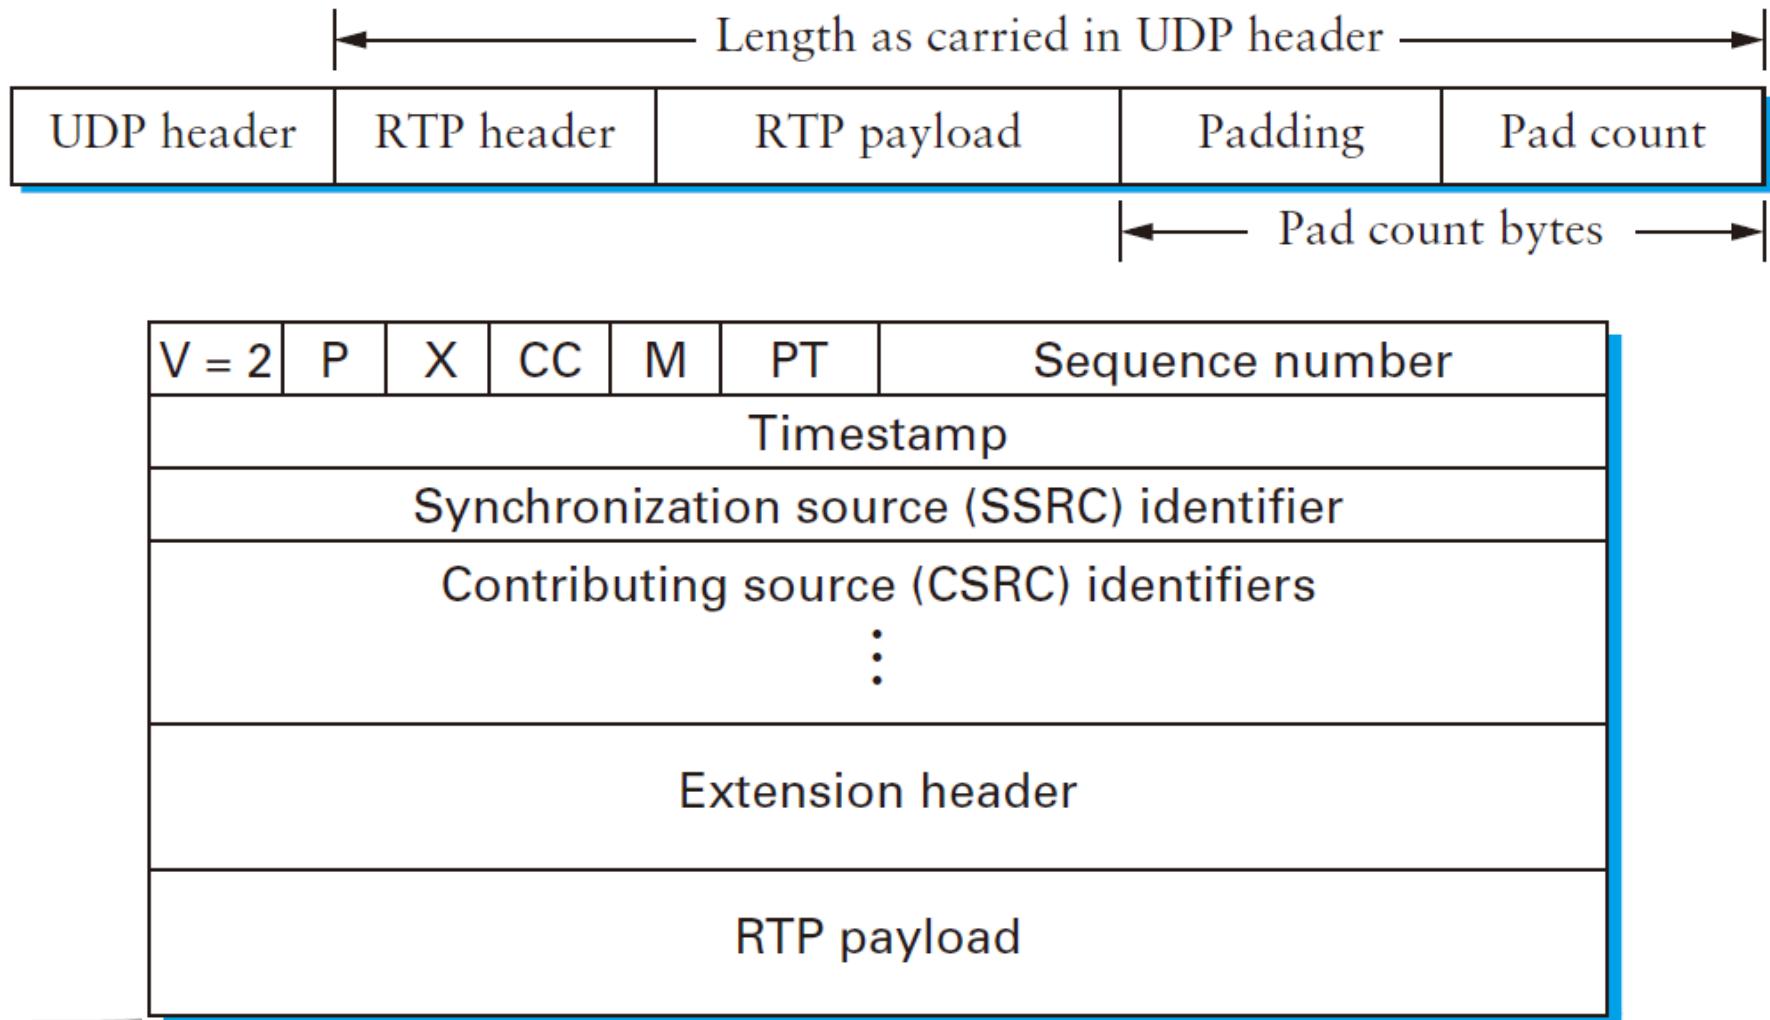


RTP and QoS

- RTP does not provide any mechanism to ensure **timely data delivery**
- RTP encapsulation is only seen at end systems, and unseen by **intermediate routers**
- Routers make **no special effort** for RTP packets

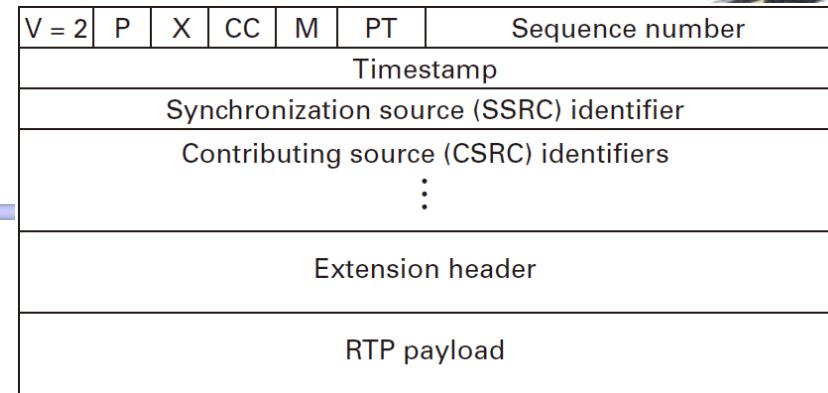


RTP Packets





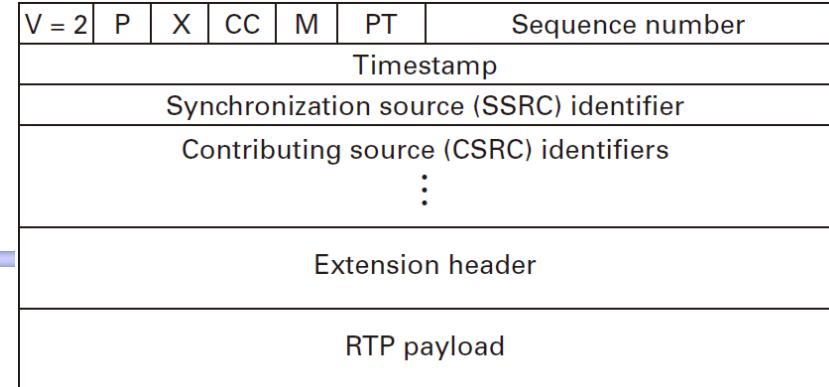
RTP Header



- Version (2 bits)
 - Version of the protocol, Current 2
- P (Padding) (1 bit)
 - Indicates if there are extra padding octets
- X (Extension) (1 bit)
 - Indicates presence of an **extension header**
- CC (CSRC Count) (4 bits)
 - Number of CSRC identifiers
- M (Marker) (1 bit)
 - Indicates special relevance for the application



RTP Header



- **Payload Type (7 bits)**

- Indicates type of encoding currently being used
- Sender can change encoding in middle of session

- **Example payload:**

- 0: PCM mu-law, 64 kbps
- 3: GSM, 13 kbps
- 7: LPC, 2.4 kbps
- 26: Motion JPEG
- 31: H.261
- 33: MPEG2 video





RTP Header

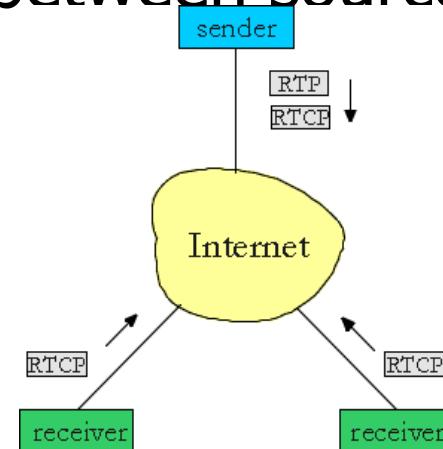
V = 2	P	X	CC	M	PT	Sequence number
						Timestamp
						Synchronization source (SSRC) identifier
						Contributing source (CSRC) identifiers
						:
						Extension header
						RTP payload

- Sequence Number (16 bits)
 - Increment by 1 for each RTP data packet sent
- Timestamp: (32 bits)
 - Sampling instant of first octet, used for play-back
- SSRC (Synchronization source identifier) (32 bits)
 - Uniquely identifies the source of a stream in a RTP session
- CSRC (Contributing source IDs) (32 bits)
 - Enumerate contributing sources to a stream
- Extension header
 - Specific header for certain payload type



RTP Control Protocol (RTCP)

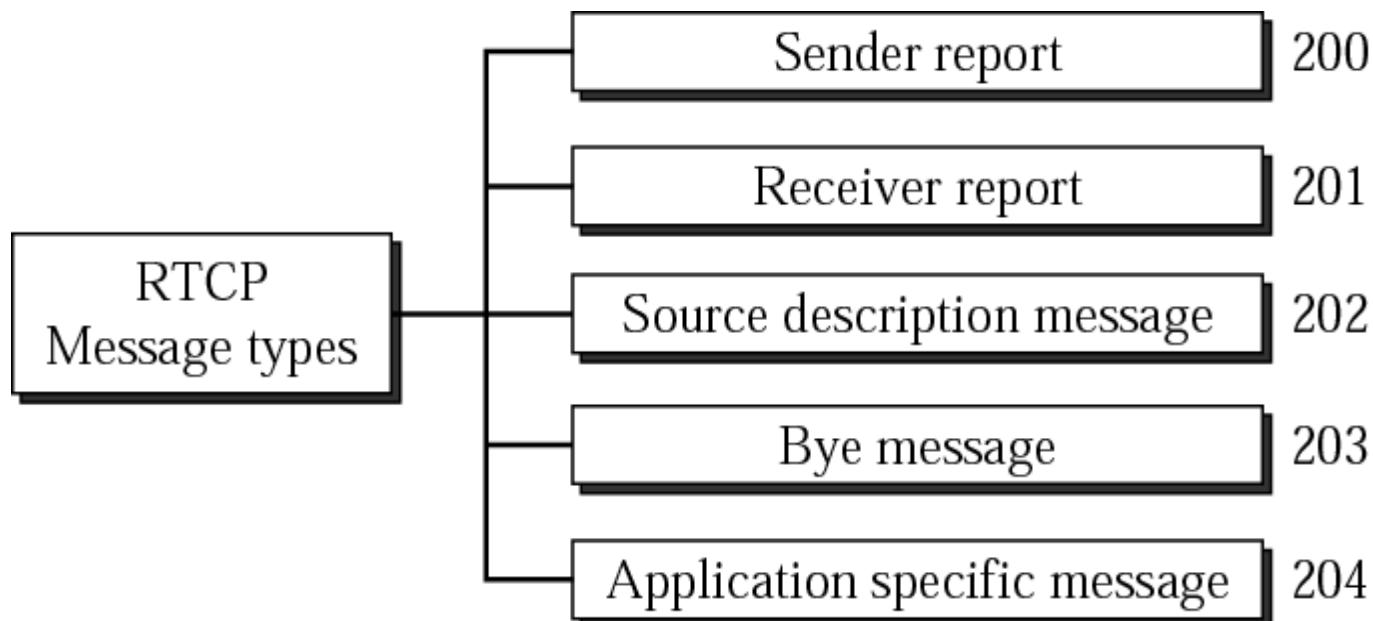
- Specifies **report PDUs** exchanged between sources and destinations
 - Receiver reception report
 - Sender report
 - Source description report
- Reports contain **statistics** such as the number of RTP-PDUs sent, number of RTP-PDUs lost, inter-arrival jitter
- Used by application to **modify sender transmission rates** and for **diagnostics** purposes





RTCP Message Types

- Several RTCP PDUs of different types can be transmitted in **a single UDP segment**





Sender/Receiver Report PDUs

V	P	RC	PT=200/201 → SR/RR	Length (16 bits)	Header						
SSRC of Sender											
NTP Timestamp, most significant word				Sender Info	Sender Info						
NTP Timestamp, least significant word											
RTP Timestamp											
Sender's PDU Count											
Sender's Octet Count											
SSRC_1 (SSRC of the 1 st Source)											
Fraction Lost	Cumulative Number of PDU Lost				Report Block 1						
Extended Highest sequence Number Received											
Interarrival Jitter											
Last SR (LSR)											
Delay Since Last SR (DLSR)											
SSRC_2 (SSRC of the 2 nd Source)											
....											
Profile-Specific Extensions											



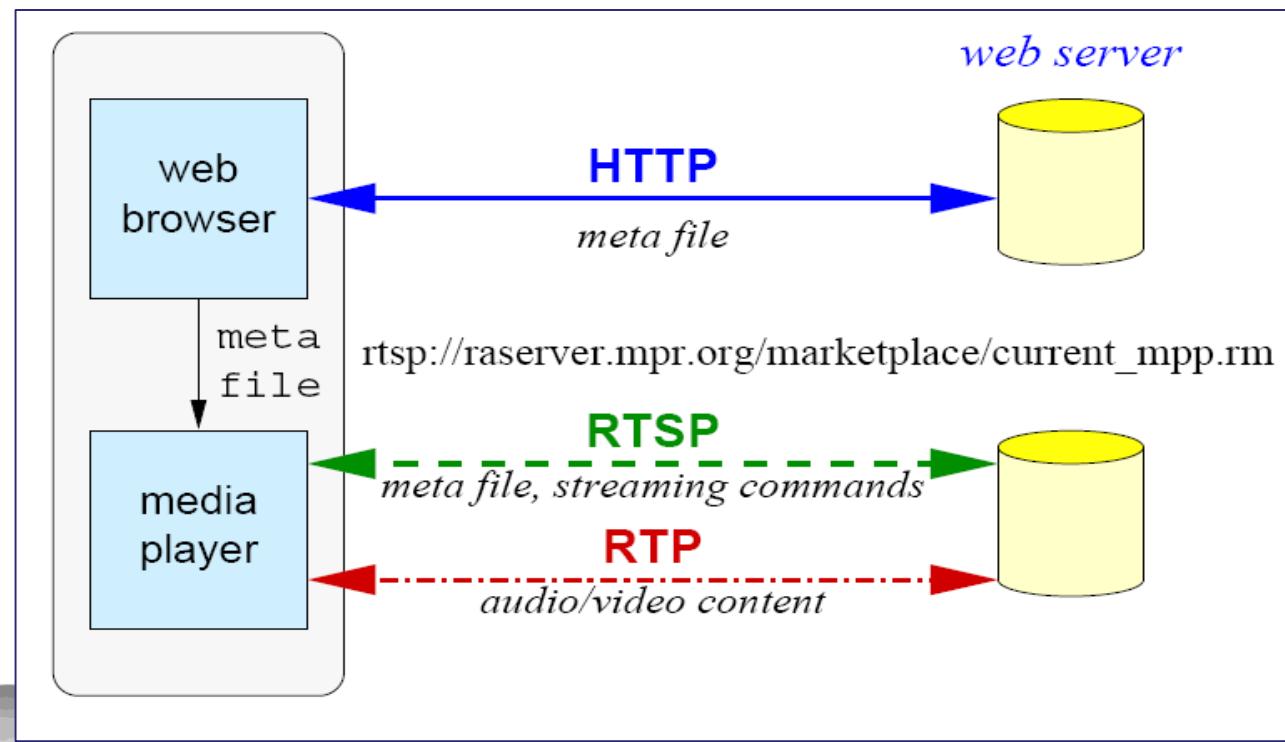
Real Time Streaming Protocol (RTSP)

- RFC 2326
 - An app-level protocol that establishes and controls media sessions between end points
 - Support VCR commands: rewind, fast forward, pause, resume, repositioning, etc...
 - Can built upon UDP or TCP, commands sent in **ASCII text**
 - Integration with web architecture, separate stream channel and control channel



RTSP Scenario

- Metafile communicated to web browser using HTTP
- Browser launches player
- Player sets up an RTSP control connection, data connection to streaming server



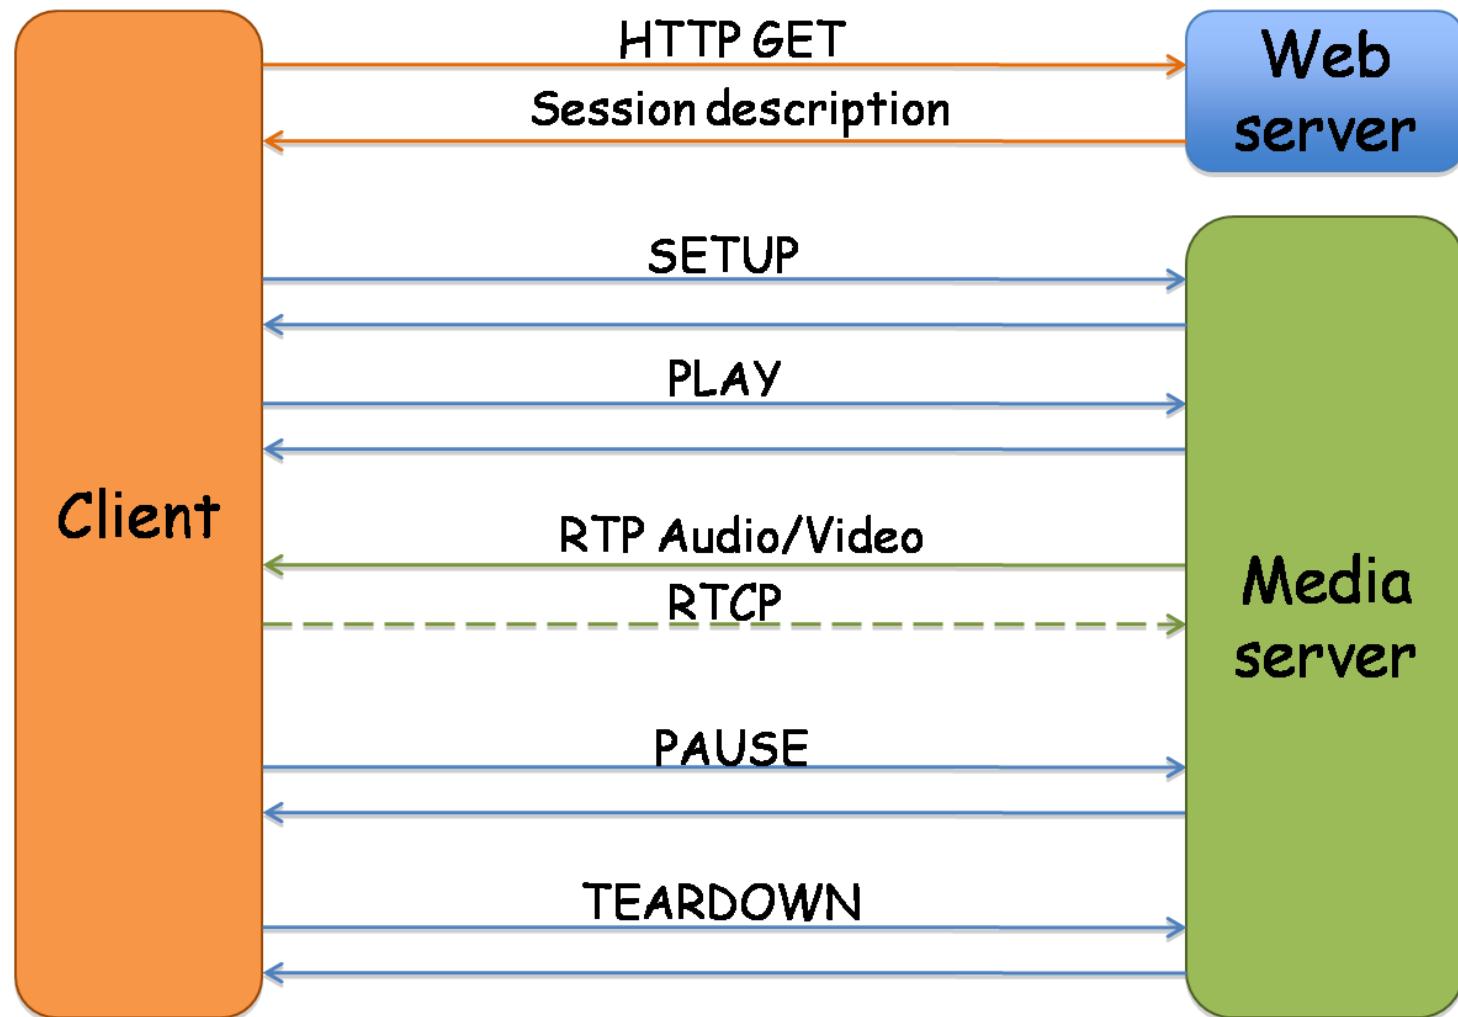


A Meta-File Example

```
<title>Twister</title>
<session>
  <group language=en lipsync>
    <switch>
      <track type=audio
        e="PCMU/8000/1"
        src = "rtsp://audio.example.com/twister/audio.en/lofi">
      <track type=audio
        e="DVI4/16000/2" pt="90 DVI4/8000/1"
        src="rtsp://audio.example.com/twister/audio.en/hifi">
    </switch>
    <track type="video/jpeg"
      src="rtsp://video.example.com/twister/video">
  </group>
</session>
```



RTSP Operation





SETUP Example

- Specifies the **transport mechanism** used for streaming media
- Client can issue SETUP to **change parameters** for already started media

Client -> Server:

SETUP rtsp://audio.example.com/twister/audio RTSP/1.0

CSeq: 302

Transport: RTP/UDP; compression; unicast; client_port=4588-4589

Server -> Client:

RTSP/1.0 200 OK

Cseq: 302

Date: 23 Jan 1997 15:35:06 GMT

Session: 47112344

Transport: RTP/UDP; compression; unicast;

client_port=4588-4589;server_port=6256-6257



PLAY Example

- Plays from beginning to end of **range** specified
- Scale header can be used to **change viewing rate**

Client -> Server:

PLAY rtsp://audio.example.com/twister.en RTSP/1.0

CSeq: 833

Session: 12345678

Range: smpte=0:10:20-;time=19970123T153600Z

Server -> Client:

RTSP/1.0 200 OK

CSeq: 833

Date: 23 Jan 1997 15:35:06 GMT

Range: smpte=0:10:22-;time=19970123T153600Z



Pause and Teardown

■ Pause

- May contain Range header to specify when to pause
- Server will terminate session after **timeout period expires**

■ Teardown

- Frees up resources on the server

Client -> Server

PAUSE rtsp://audio.example.com/twister/audio.en/lofi RTSP/1.0

CSeq: 887

Session: 12345678

Range: smpte=0:15:27

... ...

TEARDOWN rtsp://audio.example.com/twister/audio.en/lofi RTSP/1.0

CSeq: 892

Session: 12345678

Server->Client

RTSP/1.0 200 OK

CSeq: 892



RTSP Reliability

- If using TCP message is sent just once
- If using UDP, RTSP will **retransmit** if not receive ACK
 - Timeout is initially set to 500 ms
 - Can re-compute timeout based on RTT like TCP
- **Sequence no** is not incremented for retransmission
 - **Timestamp** is used to overcome retransmission ambiguity



Voice-over-IP (VoIP)

- *VoIP end-end-delay requirement:* needed to maintain “conversational” aspect
 - higher delays noticeable, impair interactivity
 - < 150 msec: good
 - > 400 msec bad
 - includes application-level (packetization, playout), network delays
- *session initialization:* how does callee advertise IP address, port number, encoding algorithms?
- *value-added services:* call forwarding, screening, recording
- *emergency services:* 911



VoIP: packet loss, delay

- ***network loss:*** IP datagram lost due to network congestion (router buffer overflow)
- ***delay loss:*** IP datagram arrives too late for playout at receiver
 - delays: processing, queueing in network; end-system (sender, receiver) delays
 - typical maximum tolerable delay: 400 ms
- ***loss tolerance:*** depending on voice encoding, loss concealment, packet loss rates between 1% and 10% can be tolerated



Session Initiation Protocol (SIP)

- RFC 3261 – an application level protocol
 - A signaling protocol used for **controlling multimedia communication sessions** such as voice and video calls over IP
 - For create, modify and terminate **two-party (unicast)** or **multiparty (multicast)** sessions consisting of one or several media streams
 - An alternative to ITU's H.323, used for IP Telephony since 1994
- Application examples
 - Video conferencing
 - Streaming multimedia distribution
 - Instant messaging
 - Online games



SIP Services

■ Setting up a call

- Determine current IP address of callee by **SIP address**
- Let callee know caller wants to establish a call
- Caller, callee agree on media type, encoding

■ Call management

- Add new media streams during call
- Change encoding during call
- Invite others, transfer and hold calls



SIP Addressing

■ Uses Internet URLs

- Uniform Resource Locators
- Supports both Internet and PSTN addresses
- General form is user@host

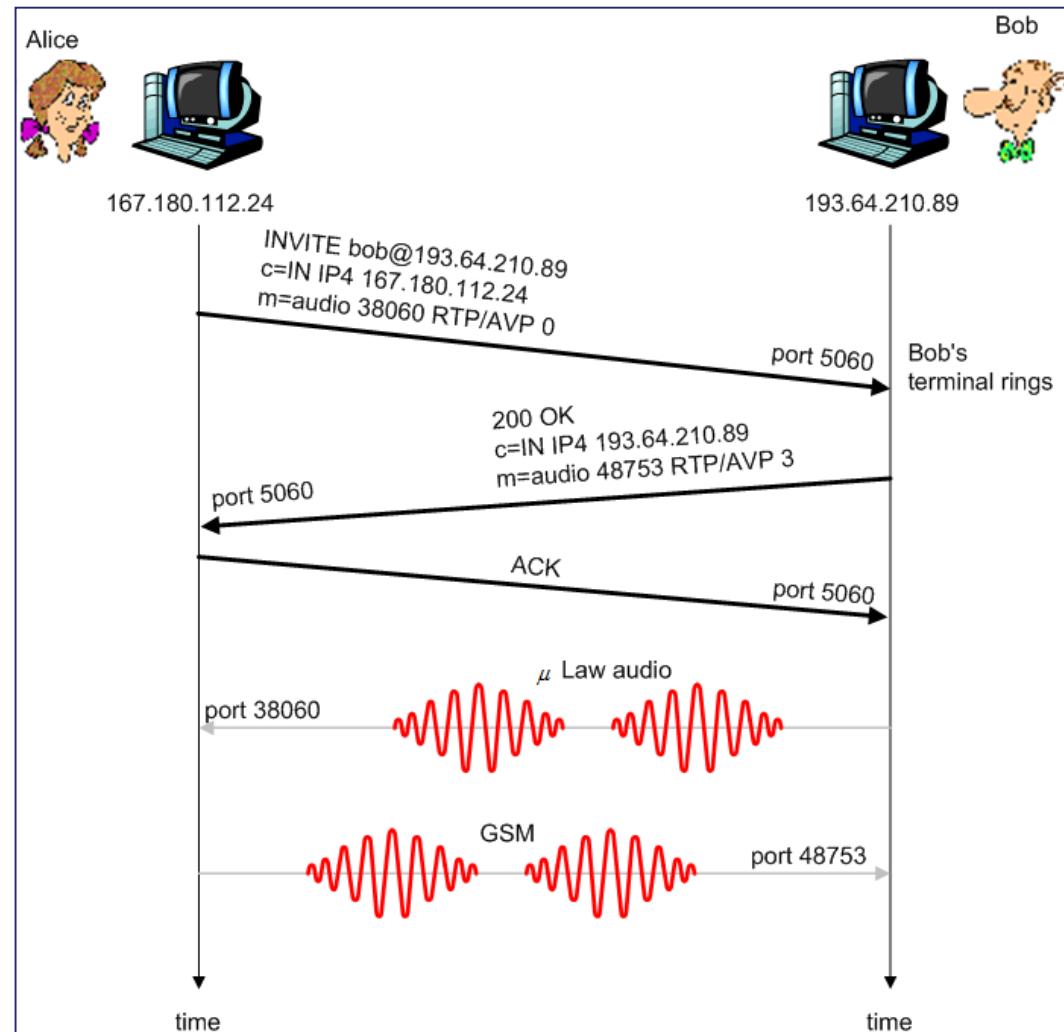
■ Examples

- sip:alan@wcom.com
- sip:J.T. Kirk <kirk@starfleet.gov>
- sip:+1-613-555-1212@wcom.com;user=phone
- sip:guest@10.64.1.1
- sip:790-7360@wcom.com;phone-context=VNET



Setting up a Call

- Alice's SIP agent invite msg, indicates her **port, IP address, encoding** she prefers to receive (PCM μ law)
- Bob's agent 200 OK msg, indicates his **port, IP address, preferred encoding** (GSM)
- SIP msgs can be **sent over TCP or UDP**, here sent over RTP/UDP
- Default SIP port number is 5060





Other Possible Reply

- **Rejecting a call**
 - Bob' agent rejects with "600 busy", "503 unavailable", "302 gone", "401 unauthorized"
- **Further negotiation**
 - Bob replies with "606 Not Acceptable", listing his encoders
 - Alice can then send new "INVITE" message, advertising different encoder



A SIP Request Message

```
INVITE sip:bob@domain.com SIP/2.0
Via: SIP/2.0/UDP 167.180.112.24:5060
From: sip:alice@hereway.com
To: sip:bob@domain.com
Call-ID: a2e3a@pigeon.here way.com
CSeq: 1 INVITE
Content-Type: application/sdp
Content-Length: 885
...
c=IN IP4 167.180.112.24
m=audio 38060 RTP/AVP 0
```

- Use HTTP message syntax
- **Via**: Shows route taken by request
- **Call-ID**: unique identifier generated by client
- **CSeq**: Command Sequence number, incremented for each successive request



A SIP Response Message

```
SIP/2.0 200 OK
Via: SIP/2.0/UDP 167.180.112.24:5060
From: sip:alice@hereway.com
To: sip:bob@domain.com
Call-ID: a2e3a@pigeon.here way.com
CSeq: 1 INVITE
```

- **Via, From, To, Call-ID, and CSeq are copied exactly from Request**
 - To and From are NOT swapped



Finding a Callee

- **SIP address** must be transformed to **IP address** of callee's current host
 - Bob may move around, getting different IP addresses (using mobile devices)
- Different SIP servers to handle this
 - **Registrar**: Accepts REGISTER requests from clients (caller or callee)
 - **Redirect**: Sends address of next hop towards callee back to caller
 - **Proxy**: Decides next hop and forwards request (as a broker)



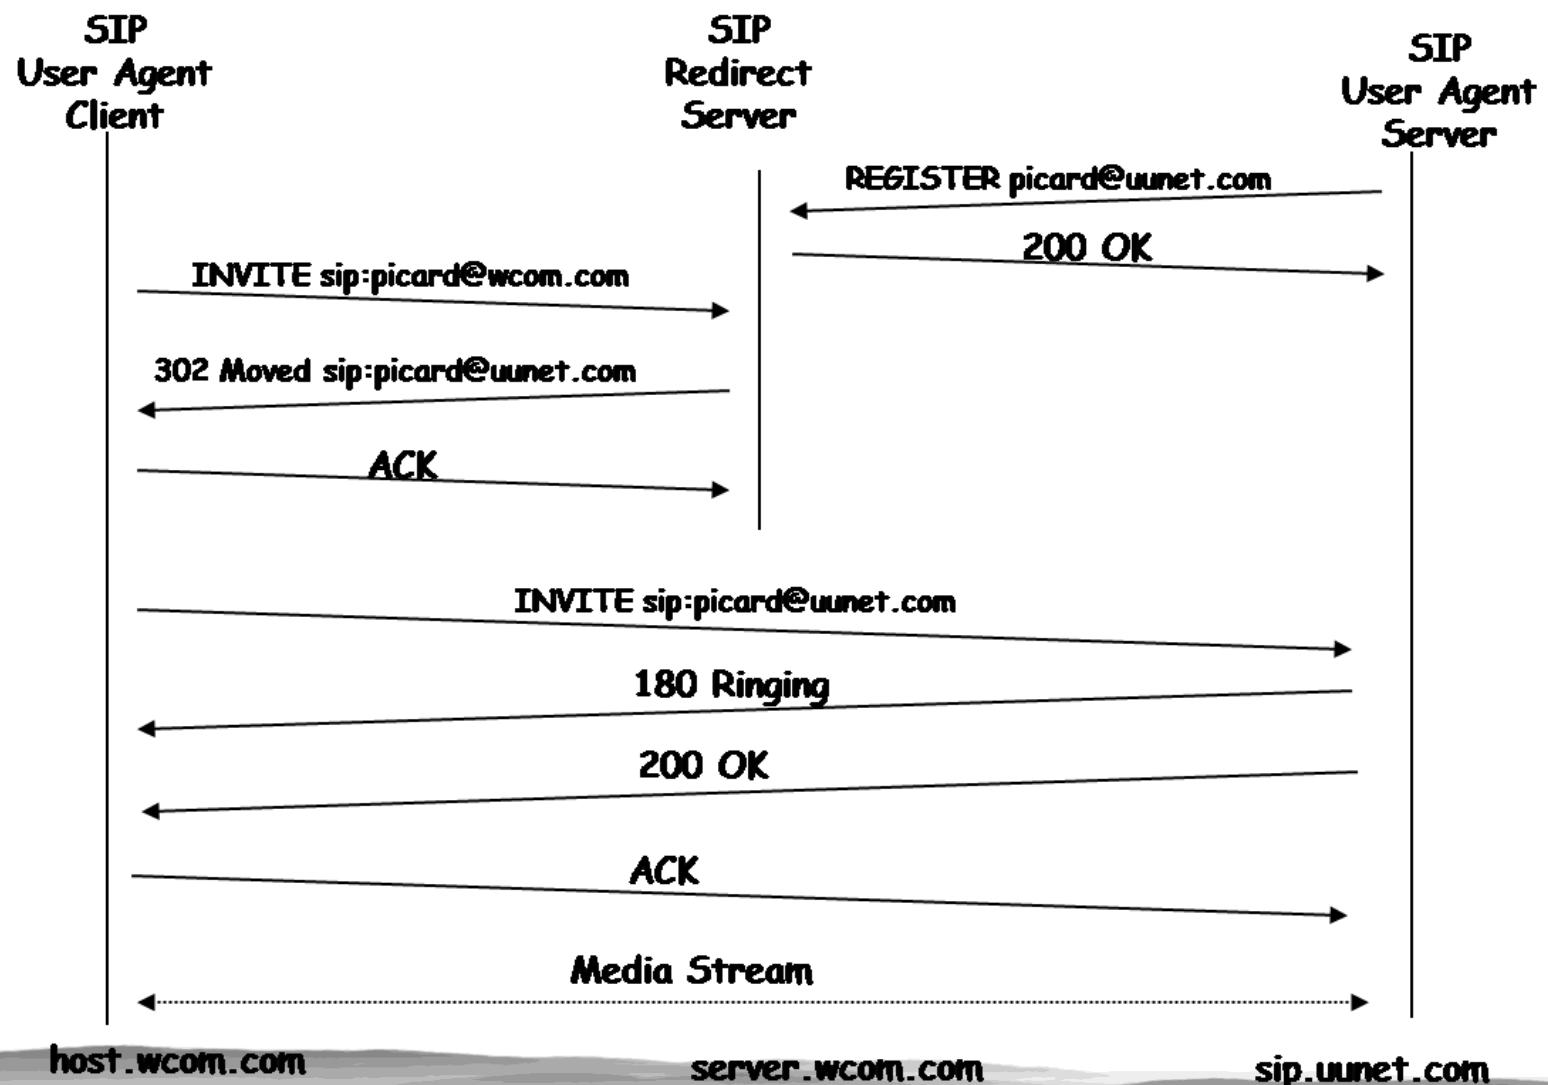
SIP Registrar

- When Bob starts SIP agent, agent sends SIP REGISTER message to Bob's registrar server
- Similar function needed by Instant Messaging

```
REGISTER sip:domain.com SIP/2.0
Via: SIP/2.0/UDP 193.64.210.89
From: sip:bob@domain.com
To: sip:bob@domain.com
Expires: 3600
```

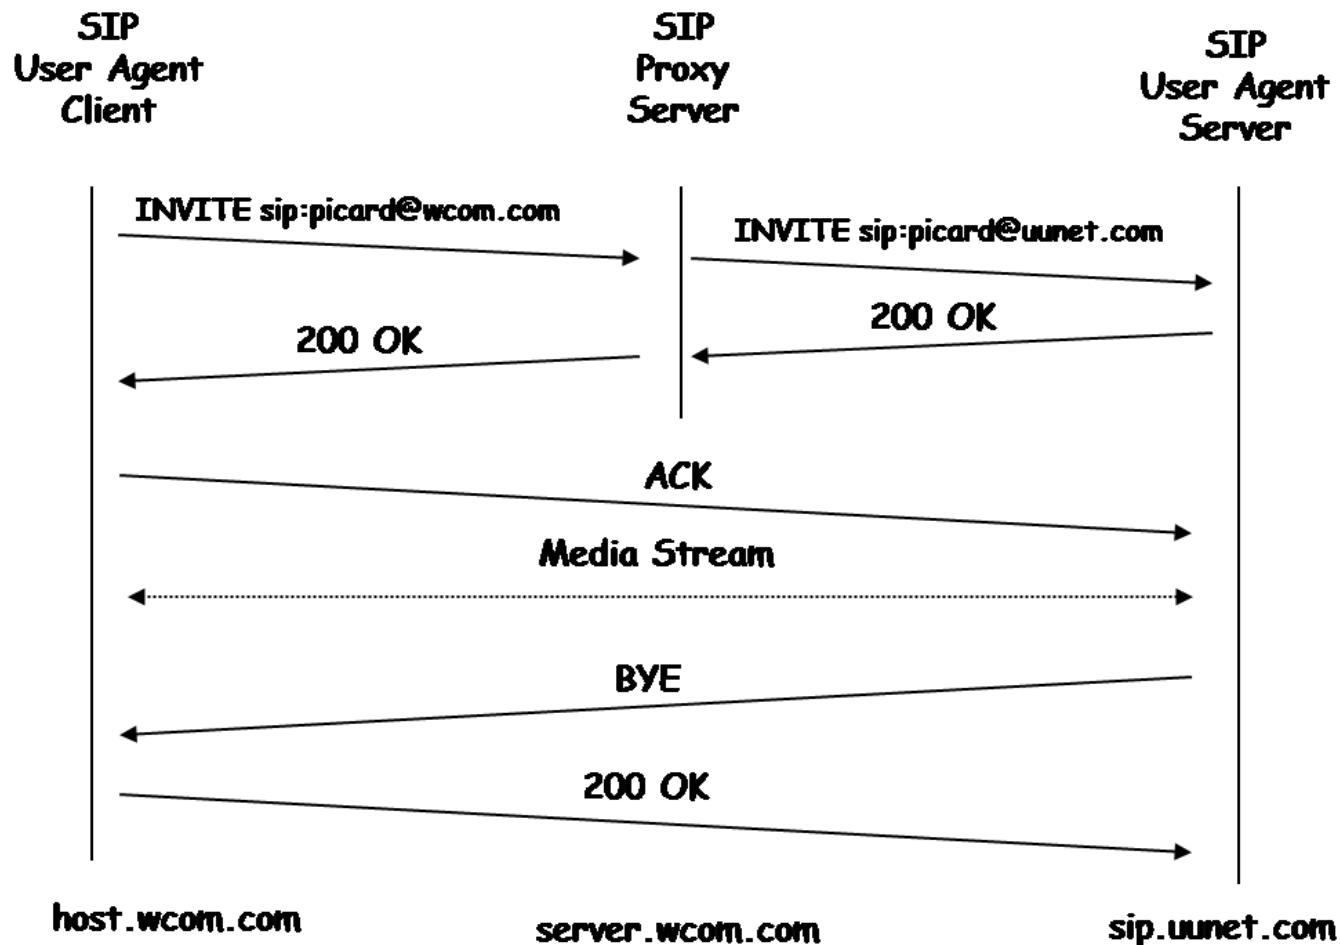


Using Redirect Server





Using Proxy Server



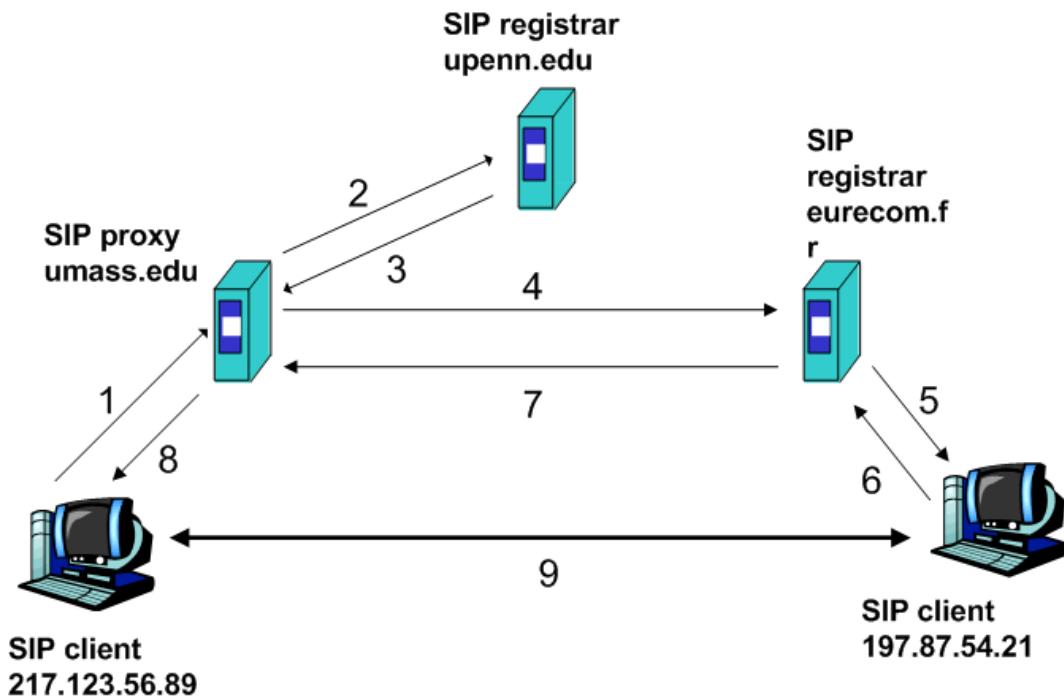


A More Complicated Example

Caller: Bob@umass.edu

Callee: Alice@upenn.edu

- (1) Bob's agent sends INVITE message to umass SIP proxy
- (2) Proxy forwards request to upenn registrar server
- (3) upenn server returns redirect response, indicating that it should try Alice@eurecom.fr
- (4) umass proxy sends INVITE to eurecom registrar
- (5) eurecom registrar forwards INVITE to 197.87.54.21, which is running Alice's SIP agent
- (6-8) SIP response sent back
- (9) media sent directly between SIP agents





Summary

- Multimedia networking
- Multimedia applications
 - Streaming, stored audio, video
 - Conversational voice/video over IP
 - Streaming live audio, video
- Protocols
 - Real-time Transport Protocol (RTP)
 - Session Initiation Protocol (SIP)
 - Real Time Streaming Protocol (RTSP)



Homework

- 第7章： P1, P3, P5